

ReadMe for CS-433 Project 1

David Howard Neill

Nathan Simonis

Ludvig Warnberg Gerdin

Contents

Summary of the contents of the files in the folder.	1
implementations.py	1
run.py	1
cross_validation.py	1
data folder	2
Description of the functions.	2
implementations.py	2
<i>Models</i>	2
<i>Cost functions</i>	5
<i>Gradients</i>	6
<i>Activation functions</i>	7
<i>Helpers</i>	7
<i>Prepare features</i>	8
<i>Performance metrics</i>	9
cross_validation.py	10

We have submitted to you a folder containing three scripts written in python as well as a sub-folder containing the data you have provided us.

Summary of the contents of the files in the folder.

implementations.py

This script contains exclusively functions that have been used throughout the project. The 6 basic methods that implement the machine learning techniques seen in the courses are included. In addition, we have included other sections where the helpers needed to apply these techniques are provided. Among others, there are sections that calculate cost functions, gradients or allow us to import data. Each section is appropriately titled and the methods that make it up are adequately commented and described.

run.py

This script uses the functions present in the previous section in order to load the data as you provided them to us, clean them, prepare the explanatory variables and perform feature augmentation, train the model, make predictions on unseen data and finally, generate a submission file that is the same as the one we used to obtain the best accuracy score on the AICrowd platform.

cross_validation.py

This script, reproduces the procedure we used to select the parameters of our model. Namely, K-fold cross-validation using a grid search algorithm to select the parameters that allow us to obtain the best average accuracy in cross-validation. Please note that we have used a fairly large grid and therefore this script takes time to complete. At the end of it, it generates a text file with the final selected parameters that were used to generate the predictions.

data folder

This folder contains .csv files. This includes, the data that you have provided us with as well as the predictions file that was used on the AICrowd platform.

Description of the functions.

implementations.py

Models

1.) `least_squares_GD(y, tx, initial_w, max_iters, gamma, verbose=False):`

Least squares with MSE loss and Gradient Descent.

Parameters

`y` : Vector

Dependent variable.

`tx` : Matrix

Explanatory variables.

`initial_w` : Vector

Initial weights.

`max_iters` : Integer scalar

Maximum number of iterations.

`gamma` : Real scalar

Learning rate.

`verbose` : Boolean, optional

Print each step of Gradient Descent or not. The default is False.

Returns

`W` : Vector

Final vector of weights.

`loss` : Real scalar

Loss given final weights.

2.) `least_squares_SGD(y, tx, initial_w, max_iters, gamma, batch_size=1, verbose=False):`

Least squares with MSE loss and Stochastic Gradient Descent.

Parameters

`y` : Vector

Dependent variable.

`tx` : Matrix

Explanatory variables.

`initial_w` : Vector

Initial weights.

`max_iters` : Integer scalar

Maximum number of iterations.

`gamma` : Real scalar

Learning rate.

`batch_size` : Integer scalar, optional

Size of the batch. The default is 1.

`verbose` : Boolean, optional

Print each step of Gradient Descent or not. The default is False.

Returns

W : Vector
 Final vector of weights.
loss : Real scalar
 Loss given final weights.

3.) least_squares(y, tx):

Linear regression fit using normal equations.

Parameters

y : Vector
 Dependent variable.
tx : Matrix
 Explanatory variables.

Returns

W : Vector
 Final vector of weights.
loss : Real scalar
 Loss given final weights.

4.) ridge_regression(y, tx, lambda_):

Ridge regression fit using normal equations.

Parameters

y : Vector
 Dependent variable.
tx : Matrix
 Explanatory variables.
lambda_ : Real scalar
 Regularization parameter.

Returns

W : Vector
 Final vector of weights.
loss : Real scalar
 Loss given final weights.

5.) logistic_regression(y, tx, initial_w, max_iters, gamma, verbose=False):

Logistic regression with log loss and Gradient Descent.

Parameters

y : Vector
 Dependent variable.
tx : Matrix

Explanatory variables.
initial_w : Vector
Initial weights.
max_iters : Integer scalar
Maximum number of iterations.
gamma : Real scalar
Learning rate.
verbose : Boolean, optional
Print each step of Gradient Descent or not. The default is False.

Returns

W : Vector
Final vector of weights.
loss : Real scalar
Loss given final weights.

6.) reg_logistic_regression(y, tx, lambda_, reg, initial_w, max_iters, gamma, verbose=False, early_stopping=True, tol=0.0001, patience=5):

Regularized logistic regression with log loss and Gradient Descent with early stopping

Parameters

y : Vector
Dependent variable..
tx : Matrix
Explanatory variables.
lambda_ : Real scalar
Regularization parameter.
reg : Integer scalar
L1 or L2 regularization.
initial_w : Vector
Initial weights.
max_iters : Integer scalar
Maximum number of iterations.
gamma : Real scalar
Learning rate.
verbose : Boolean, optional
Print each step of Gradient Descent or not. The default is False.
early_stopping : Boolean, optional
Enable early stopping or not. The default is True.
tol : Real scalar, optional
Minimum amount loss needs to change by. The default is 0.0001.
patience : TYPE, optional
Number of iterations where there must be a decrease in loss by tol. The default is 5.

Returns

W : Vector
Final vector of weights.
loss : Real scalar
Loss given final weights.

Cost functions

1.) `mse(y, tx, w):`

Mean squared error loss function.

Parameters

`y` : Vector

Dependent variable.

`tx` : Matrix

Explanatory variables.

`w` : Vector

Weights.

Returns

`loss` : Real scalar

MSE loss.

2.) `logistic_error(y, tx, w):`

Log loss function.

Parameters

`y` : Vector

Dependent variable.

`tx` : Matrix

Explanatory variables.

`w` : Vector

Weights.

Returns

`loss` : Real scalar

Log loss.

3.) `reg_logistic_error(y, tx, w, lambda_, reg):`

Log loss function with regularization term.

Parameters

`y` : Vector

Dependent variable.

`tx` : Matrix

Explanatory variables.

`w` : Vector

Weights.

`lambda_` : Real scalar

Regularization parameter

`reg` : Integer scalar

L1 or L2 regularization

Returns

`loss` : Real scalar

Log loss with regularization term.

Gradients

1.) `mse_grad(y, tx, w):`

Compute gradient for MSE loss.

Parameters

`y : Vector`

Dependent variable.

`tx : Matrix`

Explanatory variables.

`w : Vector`

Weights.

Returns

`gradient : Real scalar`

MSE gradient.

2.) `logistic_grad(y, tx, w):`

Compute gradient for log loss.

Parameters

`y : Vector`

Dependent variable.

`tx : Matrix`

Explanatory variables.

`w : Vector`

Weights.

Returns

`gradient : Real scalar`

Log loss gradient.

3.) `reg_logistic_grad(y, tx, w, lambda_, reg):`

Compute gradient for log loss with regularization term.

Parameters

`y : Vector`

Dependent variable.

`tx : Matrix`

Explanatory variables.

`w : Vector`

Weights.

`lambda_ : Real scalar`

Regularization parameter

`reg : Integer scalar`

L1 or L2 regularization

Returns

`gradient : Real scalar`

Log loss with regularization term gradient.

Activation functions

1.) sigmoid(x):

Compute sigmoid function

Parameters

x : Vector, Matrix or scalar

Explanatory variables.

Returns

sigmoid : Vector, Matrix or scalar

Sigmoid function applied to input.

Helpers

1.) batch_iter(y, tx, batch_size, num_batches=1, shuffle=True):

Generate a minibatch iterator for a dataset.

Parameters

y : Vector

Dependent variable.

tx : Matrix

Explanatory variables.

batch_size : Integer scalar

Size of the batch.

num_batches : Integer scalar, optional

Number of batches. The default is 1.

shuffle : Boolean, optional

Shuffle data or not. The default is True.

Yields

y : Vector

Mini-batch of y

tx : Matrix

Mini-batch of tx

2.) import_data(path="data/"):

Import csv files of train and test data. They must be in the same folder.

Parameters

path : String, optional

Directory of the files The default is "data/".

Returns

train : Matrix

Training set.

test : Matrix

Testing set.
col_names : Vector
Column names.

3.) create_csv_submission(ids, y_pred, name):

Creates an output file in csv format for submission to AICrowd

Parameters

ids : Vector
Event ids associated with each prediction.
y_pred : Vector
Predicted class labels.
name : String
String name of .csv output file to be created.

4.) standardize_numpy(x, mean=None, std=None):

Standardize the original data set.

Parameters

x : Matrix
Data to standardize.
mean : Vector, optional
Previously computed mean. The default is None.
std : Vector, optional
Previously computed standard deviation. The default is None.

Prepare features

1.) split_X_y(train, test, cols):

Create tx matrix for train & test + y vector for train.

Parameters

train : Matrix
Training set.
test : Matrix
Testing set.
cols : Vector
Column names.

Returns

tx_train : Matrix
tx for training set.
y_train : Vector
y for training set.
tx_test : Matrix
tx for testing set.

2.) build_poly(x, degree):

Polynomial basis functions for each column of x, for j=1 up to j=degree, and single constant term.

Parameters


```

-----
x : Matrix
    Matrix to apply augmentation on.
degree : Integer scalar
     $j^{\{th\}}$  degree for polynomial basis function.

Returns
-----
phi : Matrix
    Augmented dataset.

3.) prepare_features(tx_nan, degree, mean_nan=None, mean=None, std=None):

Clean and prepare for learning. Mean imputing, missing value indicator, standardize.
Parameters
-----
tx_nan : Matrix
    Explanatory variables.
degree : Integer scalar
     $j^{\{th\}}$  degree for polynomial basis function.
mean_nan : Vector, optional
    Compute column means with if necessary, The default is None.
mean : Vector, optional
    Previously computed mean. The default is None.
std : Vector, optional
    Previously computed standard deviation. The default is None.

Returns
-----
tx : Matrix
    Explanatory variables.
mean : Vector
    Mean of columns.
std : Vector
    Standard deviation of columns.
mean_nan : Vector
    Mean for columns with nan.
nan_cols : Vector
    Nan indicator columns.

```

Performance metrics

```

1.) logistic_prediction(tx, w):

Make a prediction with logistic regression model.
Parameters
-----
tx : Matrix
    Explanatory variables.
w : Vector
    Weights.

Returns
-----
y_pred : Vector

```

Predictions.

2.) regression_prediction(tx, w):

Make a prediction with linear regression model.

Parameters

tx : Matrix

Explanatory variables.

w : Vector

Weights.

Returns

y_pred : Vector

Predictions.

3.) f1_score(y_targ, y_pred):

Compute the F1 score of a prediction.

Parameters

y_targ : Vector

Dependent variable.

y_pred : Vector

Prediction.

Returns

score : Real scalar

Score.

4.) accuracy(y_targ, y_pred):

Compute the accuracy of a prediction.

Parameters

y_targ : Vector

Dependent variable.

y_pred : Vector

Prediction.

Returns

score : Real scalar

Score.

cross_validation.py

1.) cross_validation(y_tr, tx_tr, y_te, tx_te, comb, verbose=2):

Train model, compute in-sample and out-of-sample loss and accuracy.

Parameters

y_tr : Vector

Dependent variable in training set.

tx_tr : Matrix
 Explanatory variables in training set.
 y_te : Vector
 Dependent variable in testing set.
 tx_te : Matrix
 Explanatory variables in testing set.
 comb : Dictionnary
 Combination of parameters.
 Example : {"gamma":0.1, "lambda":0.01, "reg":2}
 verbose : Boolean, optional
 Print progress or not. The default is 2.

Returns

loss_tr : Real scalar
 Training loss.
 loss_te : Real scalar
 Testing loss
 f1 : Real scalar
 Testing F1 score.
 acc : Real scalar
 Testing accuracy.

2.) model_selection(y, tx, k_fold, degree, grid, seed, verbose=2):

Select the best model from all possible combinations of grid.

Parameters

y : Vector
 Dependent variable.
 tx : Matrix
 Explanatory variables.
 k_fold : Integer scalar
 Number of folds for cross-validation
 degree : Integer scalar
 j^{th} degree for polynomial basis function.
 grid : Dictionnary
 Set of all possible combinations.
 seed : Integer scalar
 Random seed.
 verbose : Boolean, optional
 Print progress or not. The default is 2.

Returns

params : Dictionnary
 Best parameters for given grid.

3.) build_k_indices(y, k_fold, seed):

Build k indices for k-fold.

Parameters

y : Vector

Dependent variable.
k_fold : Integer scalar
Number of folds for cross-validation.
seed : Integer scalar
Random seed.

Returns

k_indices : Matrix
k indices for K-fold.

4.) prepare_split_data(y, tx, degree, k_fold, seed):

Split the dataset based on k-fold cross validation and prepare features.

Parameters

y : Vector
Dependent variable.
tx : Matrix
Explanatory variables.
degree : Integer scalar
 j^{th} degree for polynomial basis function.
k_fold : Integer scalar
Number of folds for cross-validation.
seed : Integer scalar
Random seed.

Returns

y_trs : Vector
Dependent variable for training set.
tx_trs : Matrix
Explanatory variables for training set.
y_tes : Vector
Dependent variable for testing set.
tx_tes : Matrix
Explanatory variables for testing set.